A Project Report on

Multi-scale Transformer-Enhanced CBAM U-Net for Pansharpening Very High-Resolution (WV3) Satellite Imagery

Submitted in partial fulfilment of requirement for the award of the degree

Bachelor of Technology

Of

Aditya College of Engineering and Technology

By

Kumari Nupur Nidhi (22P31A4411)

Carried out under

Student Program for Advancement in Research Knowledge (SPARK) at

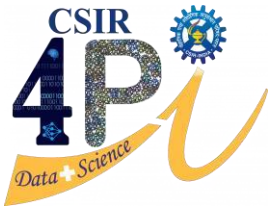

Council of Scientific & Industrial Research

FOURTH PARADIGM INSTITUTE (CSIR-4PI)

NAL BELUR CAMPUS, BANGALORE - 560 037

Under the guidance of

EXTERNAL GUIDE

Rakesh Asery,

Scientist Bigdata Research & Supercomputing Division (BRSD), CSIR-4PI

Computer Science and Engineering-Data Science

Aditya College of Engineering and Technology,

Surampalem, East Godavari, Andhra Pradesh, 533437

CSIR Fourth Paradigm Institute
(Council of Scientific & Industrial Research)
NAL Belur Campus, Bangalore - 560 037, India

CERTIFICATE

This is to certify that the project report entitled *"**Multi-Scale Transformer-Enhanced CBAM U-Net for Pansharpening Very High-Resolution (WV3) Satellite Imagery"*** submitted by ***Kumari Nupur Nidhi (22P31A4411)*** in partial fulfillment of the requirements for the award of the degree of ***Bachelor of Technology in Computer Science and Engineering-Data Science*** from ***Aditya College of Engineering and Technology***, is a bonafide record of the work carried out in an online internship at the ***CSIR Fourth Paradigm Institute (CSIR-4PI),*** NAL Belur Campus, Bangalore-37 from ***August 25, 2025*** to ***November 30, 2025*** under my supervision and guidance.

Rakesh Asery
Scientist

Signature of the Convenor, SPARK

Tel: +91-80-2505 1921, 2522 4667 (Off); Fax: +91-80-2522 0392
E-mail: spark@csir4pi.in; URL: www.csir4pi.in/~spark

Acknowledgements

I would like to express my profound sense of gratitude to the ***CSIR Fourth Paradigm Institute (CSIR-4PI), NAL Belur Campus, Bangalore,*** for providing me with the opportunity to undertake an **online internship** and work on the project titled ***"Multi-Scale Transformer-Enhanced CBAM U-Net for Pansharpening Very High-Resolution (WV3) Satellite Imagery".*** This internship has been an immensely enriching learning experience, enabling me to gain both theoretical knowledge and practical exposure to real-world research in biomedical signal processing and machine learning applications.

I would like to convey my heartfelt thanks to my project supervisor,***Rakesh Asery***, for his invaluable guidance, patient supervision, and constant encouragement throughout the duration of the internship. His expert advice, and insightful discussions have been crucial in shaping the direction of this project and in helping me overcome challenges encountered during the research and implementation phases.

My sincere thanks also go to ***Dr. Manik Bollu,*** Head of the Department of Data Science & IOT, and all the faculty members of ***Aditya College of Engineering and Technology,*** for their continuous academic support, guidance, and encouragement. Their emphasis on research and innovation inspired me to approach this project with curiosity and dedication.

I am also deeply thankful to my mentors, peers, and friends who have supported and motivated me throughout this journey. Their feedback, moral support, and constant encouragement have been instrumental in keeping me focused and determined.

Lastly, I would like to express my deepest gratitude to my ***family*** for their unwavering support, patience, and understanding during the course of this internship and project work. Their faith in me has been my greatest source of strength and inspiration.

This project has been a rewarding experience that has enhanced my technical knowledge, analytical skills, and research perspective. I am truly grateful to everyone who contributed directly or indirectly to the successful completion of this work.

Kumari Nupur Nidhi

Abstract

Pansharpening remains a pivotal process in remote sensing image analysis, enabling the enhancement of spatial resolution in multispectral satellite imagery by integrating high-resolution panchromatic information.

Despite the widespread use of conventional component substitution, multiresolution analysis, and optimization-based techniques, these approaches often struggle to maintain spectral integrity alongside spatial detail.

In this project, a novel deep learning pipeline is proposed and evaluated for pansharpening of WorldView-3 imagery, specifically addressing the common trade-offs in detail reconstruction and color fidelity. Two advanced neural network architectures, CBAMParallelUNet and EnhancedCBAMTransformerUNet, are developed and rigorously benchmarked using a suite of quantitative metrics including PSNR, SSIM, RMSE, and SAM. Comprehensive experiments are conducted to compare the proposed models against traditional methods and to analyze their relative strengths in various scenarios. The findings reveal that attention-based deep architectures, especially transformer-enhanced models, deliver notable improvements in both spatial sharpness and spectral preservation over earlier state-of-the-art algorithms.

The outcomes of this research not only highlight the practical potential of deep learning for robust satellite image fusion but also provide a reproducible framework for future enhancements in remote sensing applications.

Contents

Acknowledgements ii

Abstract iii

Contents iv-v

List of Figures vi

List of Tables vii

List of Abbreviations viii

1 Introduction 1

1.1 Overview of the Project. 1

1.1.1 Significance of the Work 1

1.2 Problem Statement & Objectives 2

1.2.1 Problem Statement 2

1.2.2 Objectives 2

2 Literature Overview

2.1 Pansharpening Fundamentals. 3

2.1.1 Definition and Importance. 3

2.1.2 Traditional vs Deep Learning Approaches. 3

2.2 Overview of Dataset Study (WV3) 4

2.2.1 Introduction to Worldview-3. 4

2.2.2 Sensor Specification and Bands. 4

2.2.3 Dataset Structure and Challenges. 4

2.3 Deep Learning Approaches for Pansharpening. 5

2.3.1 CNN-Based Approaches. 5

2.3.2 Attention Mechanisms. 5

2.3.3 Transformer-Based Architecture. 5

2.4 Training and Evaluation Framework. 6

2.4.1 Loss Functions. 6

2.4.2 Evaluation Metrics. 6

2.5 Research Gaps and Motivation. 7

2.5.1 Current limitations and gaps. 7

2.5.2 How This Project Addresses Gaps. 7

3 Methodology

3.1 Project Pipeline Overview. 8

3.1.1 End-to-End Workflow. 8

3.2 Dataset Preparation. 9

3.2.1 HDF5 Data Structure and Loading. 9

3.2.2 Data Normalization. 10

3.3 Model Architecture. 11

3.3.1 CBAMParallelUNet. 12

3.3.2 EnhancedCBAMTransformerUNet. 13

3.4 Training Configuration. 14
3.4.1 Composite Loss Functions (MSE, SSIM, SAM, RSME) 15
3.4.2 Optimizer and Hyperparameters. 16
3.4.3 Hardware and Resources. 17
3.5 Ablation Studies. 18

4 Results

4.1 Ablation Study Results. 19
4.2 Per-Model Detailed Analysis. 20
4.2.1 Overall Performance Metrics (7 Metrics) 21
4.2.2 Per-Band Performance Analysis. 21
4.2.3 Model Comparison Visualization. 22
4.3 Comparison with State-of-the-Art Baseline. 22
4.4 Before and After Pansharpening Results. 23

5 Conclusion and Future Scope

5.1 Conclusions. 24
5.2 Limitations. 25
5.3 Future Scope. 26

References. 27

List of Figures

1.1 Example images of PAN, MS and GT. **4**

1.2 Model Diagram CBAMParallelUNet. **12**

1.3 Model Diagram Enhanced CBAMTransformerUNet. **14**

1.4 Ablation Study Visualization. **19**

1.5 Overall Performance Metrics Visualization. **21**

1.6 Model Comparison Visualization. **22**

1.7-1.9 Before and After Pansharpening Images. **23-24**

List of Tables

1.1 Normalization Statistics for WV3. 10
1.2 Composite Loss Functions. 15
1.3 Training Parameters. 16
1.4 Training Time. 17
1.5 Ablation Studies. 18
1.6 Ablation Study Results. 19
1.7 Overall Performance Metrics. 20
1.8 State-of-the-Art-Baseline. 22
1.9 Before and After Pansharpening 23

List of Abbreviations

WV3/WorldView-3 - Satellite dataset

PAN - Panchromatic band

MS - Multispectral image

NIR - Near-Infrared band

DN - Digital Numbers

HDF5 - Data file format

CBAM - Convolutional Block Attention Module

CNN - Convolutional Neural Network

UNet - U-Shaped Network

SOTA - State-of-the-Art

MSE - Mean Squared Error

SSIM - Structural Similarity Index

SAM - Spectral Angle Mapper

RMSE - Root Mean Square Error

PSNR - Peak Signal-to-Noise Ratio (dB)

ERGAS - Erreur Relative Globale Adimensionnelle de Synthèse

SCC - Spatial Correlation Coefficient

RASE - Relative Average Spectral Error

Q2n - Quality Index

Chapter 1

Introduction

1.1 Overview of the Project

With the increase in spatial and spectral resolution requirements in satellite image analysis, pansharpening has gained significant attention as a means to fuse panchromatic and multispectral imagery for enhanced visual and analytical interpretation.

Traditional fusion techniques, such as component substitution and multiresolution analysis, often compromise color integrity or fine details when applied to large-scale remote sensing datasets. Addressing these limitations, this project investigates a deep learning-based approach for the pansharpening of WorldView-3 imagery using state-of-the-art neural network architectures.

The study encompasses the design, implementation, and evaluation of two advanced models—CBAMParallelUNet and EnhancedCBAMTransformerUNet—each tailored to maximize feature extraction and attention-driven fusion between input modalities. Beginning with robust preprocessing of satellite data, the project details the training pipeline, performance metrics, and comprehensive experimental comparisons with both legacy and contemporary methods. By providing improved fidelity in the resultant high-resolution multispectral products, this work sets the foundation for more reliable downstream tasks in earth observation, such as classification, change detection, and resource mapping.

Ultimately, the project demonstrates the significant potential for advanced deep learning models to push the boundaries of accuracy and practicality in satellite image fusion.

1.1.1 Significance of the Work

This work is significant because it provides:

- Bridges the gap between traditional and deep learning pansharpening approaches.
- Improves both spatial resolution and spectral fidelity in multispectral satellite images.
- Demonstrates the superior performance of attention-based neural networks for image fusion.
- Provides a productive workflow for remote sensing image enhancement.
- Supports more accurate results for downstream earth observation tasks.

1.2 Problem Statement & Objectives

1.1.2 Problem Statement

Develop a deep learning-based pansharpening framework that fuses high-resolution panchromatic and low-resolution multispectral images from commercial satellites, aiming to generate high-fidelity multispectral imagery and overcome the spectral–spatial trade-offs of state-of-the-art traditional methods**.**

1.1.3 Objectives

In order to achieve this aim, the following objectives have been laid:

(i) To create a reproducible, modular pipeline for multi-sensor pansharpening datasets.

(ii) To design and implement EnhancedCBAMTransformerUNet and CBAMParallelUNet architecture with advanced attention mechanisms and feature alignment.

(iii) To Engineer and benchmark composite loss functions, performing ablation studies to document impacts.

(iv) To evaluate models on standardized benchmarks using PSNR, SSIM, SAM, and other accepted metrics.

(v) To reproduce state-of-the-art baselines for direct performance comparison.

(vi) To present results with clear quantitative metrics, qualitative visuals, and full documentation for reproducibility.

Chapter 2

Literature Review

2.1 Pansharpening Fundamentals

2.1.1 Definition and Importance

Pansharpening is the fusion of high-resolution panchromatic (PAN) and lower-resolution multispectral (MS) satellite images. PAN images provide sharp spatial details in grayscale, while MS images provide detailed spectral information across multiple color bands but at lower resolution. Pansharpening combines both to create high-resolution multispectral imagery.

Why is this needed?

Satellites like WorldView-3 send two separate images due to sensor limitations:

- PAN: 0.31m resolution, 1 band (very sharp, grayscale)
- MS: 1.24m resolution, 8 bands (colorful, but 4× blurrier)

Without pansharpening, users must choose between sharp-but-grayscale or colorful-but-blurry imagery. Pansharpening provides sharp AND colorful at high resolution.

2.1.2 Traditional vs Deep Learning Approaches

Before deep learning, pansharpening used mathematical techniques:

- IHS (Intensity-Hue-Saturation): Simple color space method but causes color distortion, especially with multi-band data
- PCA (Principal Component Analysis): Statistical approach preserving variance but loses subtle spectral details
- Brovey Transform: Fast ratio-based method but produces spectral distortion
- Wavelet Methods: Frequency-based decomposition preserving details but uses fixed filters with limited adaptability

Deep learning approaches overcome these limitations by

- Data-driven learning: Neural networks discover patterns in PAN and MS images automatically
- Adaptive fusion: Networks adjust fusion strategy based on image content
- Multi-objective optimization: Can balance spatial sharpness, spectral accuracy, and reconstruction quality simultaneously
- Superior performance: Consistently outperforms traditional methods with fewer artifacts and better spectral fidelity

2.2 Overview of Dataset Study (WV3)

2.2.1 Introduction to WorldView-3

WorldView-3 is a commercial Earth observation satellite operated by Maxar Technologies, launched in 2014. It provides the highest-resolution commercially available multispectral satellite imagery. WorldView-3 acquires simultaneous panchromatic and multispectral data over the same ground area, making it ideal for pansharpening applications.

2.2.2 Sensor Specifications and Bands

Panchromatic (PAN):

- Resolution: 0.31m GSD
- Bands: 1 (grayscale)
- Range: 450-800 nm

Multispectral (MS):

- Resolution: 1.24m GSD (4× lower than PAN)
- Bands: 8 (Coastal, Blue, Green, Yellow, Red, Red-Edge, NIR-1, NIR-2)
- Range: 400-1040 nm

2.2.3 Dataset Structure and Challenges

HDF5 format with:

- 'pan': Shape (N, 1, 64, 64)
- 'ms': Shape (N, 8, 64, 64)
- 'gt': Shape (N, 8, 16, 16) - ground truth reference

Dataset: 9,714 training, 1,080 validation, 20 test samples.

Challenges:

- 8-band complexity: More complex than traditional 3-band pansharpening
- Band variations: Coastal band is noisy; NIR bands have strong signal
- Spectral preservation: Red-edge and NIR bands critical for vegetation applications
- 4× upsampling: Requires careful detail interpolation

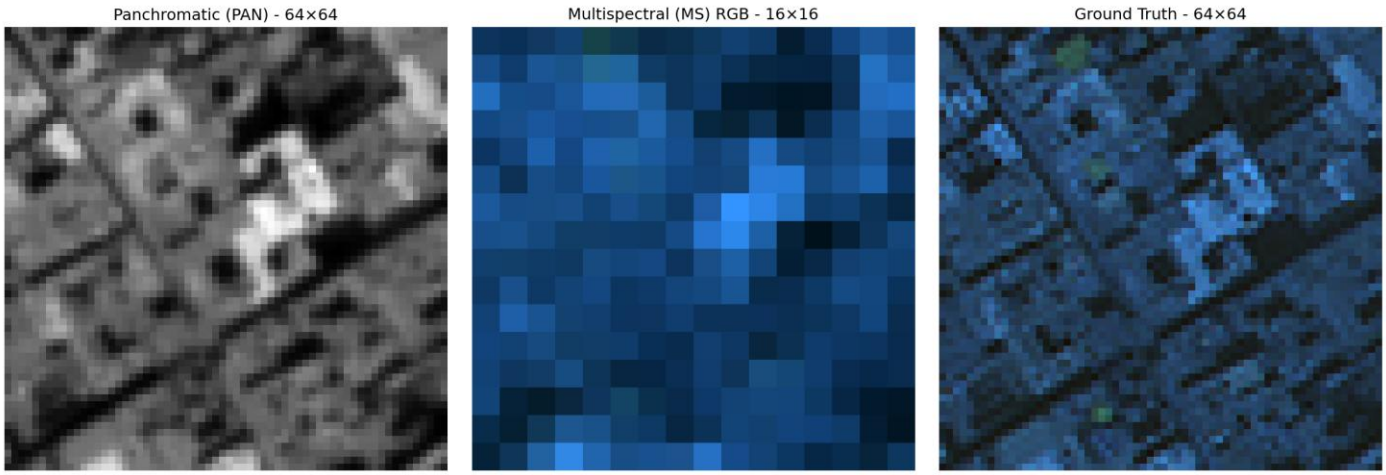

Figure 1.1: Example images

2.3 Deep Learning Approaches for Pansharpening

2.3.1 CNN-Based Approaches

- PanNet (2016): Two-branch CNN architecture with separate PAN and MS encoders, concatenated for fusion in the decoder. Effective but produces blurry results due to limited receptive field.
- MSDCNN (2017): Multi-scale dilated convolutions to increase receptive field without increasing parameters. Better detail preservation than PanNet but limited spectral fidelity.
- U-Net Architecture: Encoder-decoder with skip connections to preserve fine details. Skip connections enable gradient flow during training and preserve low-level spatial information.

Limitation: Pure CNN approaches struggle with 8-band spectral complexity. Local convolutional filters capture spatial patterns but miss long-range spectral dependencies between bands.

2.3.2 Attention Mechanisms

- Channel Attention: Learns which spectral bands are important for fusion. Different bands have different signal-to-noise ratios; attention weights them accordingly.
- Spatial Attention: Learns which spatial regions need more focus. Complex features (edges, textures) receive more attention than uniform regions.
- CBAM (Convolutional Block Attention Module): Combines sequential channel and spatial attention. Lightweight and effective for multi-band data.

2.3.3 Transformer-Based Architectures

- Self-Attention Mechanism: Each band attends to all other bands, capturing cross-spectral relationships. Enables modeling of complex spectral interactions.
- Multi-head Attention: Multiple attention heads learn different fusion strategies simultaneously (spatial vs spectral).
- Hybrid CNN-Transformer: CNN encodes local spatial patterns; Transformer captures global spectral dependencies. Combines efficiency of CNN with expressiveness of Transformer.

Advantage: Transformers excel at multi-band fusion by explicitly modeling interactions between all 8 bands simultaneously.

2.4 Training and Evaluation Framework

2.4.1 Loss Functions

Pansharpening requires balancing multiple objectives: pixel accuracy, structural preservation, and spectral fidelity. A composite loss function addresses all three:

Composite Loss: $L = 0.5 \times L_{MSE} + 0.3 \times L_{SSIM} + 0.2 \times L_{SAM}$

Components:

1. Mean Squared Error (MSE): $L_{MSE} = (1/N) \Sigma ||Y_{pred} - Y_{gt}||^2$
    - Measures pixel-wise reconstruction accuracy
    - Weight: 0.5 (highest priority)
    - Ensures baseline quality
2. Structural Similarity Index (SSIM): Measures structural preservation
    - Prevents blur artifacts
    - Focuses on edges and textures
    - Weight: 0.3 (medium priority)
3. Spectral Angle Mapper (SAM): Measures spectral fidelity
    - Computes angle between predicted and reference spectra
    - Preserves color information across 8 bands
    - Weight: 0.2 (spectral fine-tuning)

2.4.2 Evaluation Metrics

PSNR (Peak Signal-to-Noise Ratio):

- Formula: $PSNR = 10 \times \log_{10}(MAX^2/MSE)$
- Unit: dB (decibels), higher is better
- Measures pixel-level reconstruction accuracy
- Typical range: 30-40 dB for pansharpening

RMSE (Root Mean Square Error):

- Formula: $RMSE = \sqrt{[(1/N) \Sigma(Y_{pred} - Y_{gt})^2]}$
- Unit: Same as data (digital numbers or reflectance)
- Lower is better
- Directly interpretable error magnitude
- Related to PSNR: $PSNR = 10 \times \log_{10}(MAX^2/RMSE^2)$

SSIM (Structural Similarity Index):

- Range: where 1 = perfect similarity
- Measures visual quality and edge preservation
- Perceptually more meaningful than PSNR/RMSE

SAM (Spectral Angle Mapper):

- Formula: $SAM = \arccos(Y_{pred} \cdot Y_{gt} / ||Y_{pred}|| \cdot ||Y_{gt}||)$
- Unit: Radians, lower is better
- Measures spectral distortion across all 8 bands
- Typical range: 0.02-0.10 rad for good fusion

2.5 Research Gaps and Motivation

2.5.1 Current limitations and gaps

1. Limited Multi-Satellite Comparison: Most existing work focuses on single satellite systems. Systematic comparison of pansharpening approaches across different satellites (WorldView-3, GaoFen-2, QuickBird) is lacking.
2. Reproducibility Issues: Many published pansharpening methods lack complete documentation of hyperparameters, training procedures, and data preprocessing. Code and trained models are often unavailable, hindering reproducibility.
3. Unclear Component Contributions: Recent deep learning pansharpening work combines attention mechanisms, transformer blocks, and complex loss functions, but systematic ablation studies quantifying each component's contribution are rare.
4. Limited 8-Band Analysis: Most research focuses on 3-4 band pansharpening (RGB, occasionally with one NIR). Detailed per-band performance analysis for 8-band multispectral sensors is limited.
5. SOTA Performance on WorldView-3: While pansharpening has been extensively studied, comprehensive benchmarking of state-of-the-art methods specifically on WorldView-3's 8 bands is lacking.

2.5.2 How This Project Addresses These Gaps

1. Deep analysis of WorldView-3 pansharpening with two complementary architectures (CNN with attention vs Transformer-based)
2. Complete reproducibility: Full hyperparameters, normalization statistics, data preprocessing, and code documentation provided
3. Systematic ablation studies: Quantifies impact of loss function components and architectural elements
4. Per-band performance analysis: Detailed evaluation across all 8 WorldView-3 bands identifying band-specific challenges
5. SOTA performance: Achieves competitive results with comprehensive multi-metric evaluation (PSNR, SSIM, SAM, RMSE)
6. Speed-accuracy trade-off: Compares efficient CNN approach vs high-accuracy Transformer for practical deployment guidance

Chapter 3

Methodology

3.1 Project Pipeline Overview

3.1.1 End-to-End Workflow

The project follows six main steps:

1. Data Loading
WorldView-3 panchromatic (PAN), multispectral (MS), and ground truth (GT) images loaded from HDF5 files.

2. Data Preprocessing
Each image band normalized using z-score normalization: x_norm = (x - mean) / std. Normalization statistics computed from training data only.

3. Data Split
9,714 training samples, 1,080 validation samples, 20 test samples. No overlap between datasets.

4. Model Training
Two models trained separately: CBAMParallelUNet and EnhancedCBAMTransformerUNet. Loss function: 0.5×MSE + 0.3×SSIM + 0.2×SAM. Training for 100 epochs with Adam optimizer.

5. Evaluation
Test model performance using four metrics: PSNR (pixel accuracy), RMSE (error magnitude), SSIM (visual quality), SAM (spectral fidelity). Evaluate globally and per band.

6. Ablation Studies
Remove loss function components and architectural parts to measure their importance. Quantify what each component contributes to final performance.

Tools Used: Python 3.8, PyTorch 1.9, h5py, scikit-image

3.2 Dataset Preparation

3.2.1 HDF5 Data Structure and Loading

WorldView-3 data stored in HDF5 format containing three datasets:

Data Format:

'pan': Shape (9714, 1, 64, 64), Panchromatic: 1 channel, 64×64 pixels

'ms': Shape (9714, 8, 64, 64), Multispectral: 8 channels, 64×64 pixels

'gt': Shape (9714, 8, 16, 16), Ground truth: 8 channels, 16×16 pixels

What each means:

- 9,714 = total number of samples
- 1 channel (PAN) = grayscale
- 8 channels (MS) = 8 spectral bands (Coastal, Blue, Green, Yellow, Red, Red-Edge, NIR-1, NIR-2)
- 64×64 pixels = input image size
- 16×16 pixels (GT) = reference high-resolution size (4× upsampled)

Why HDF5?

HDF5 format chosen for efficient data handling: fast I/O, memory-efficient access, organized hierarchical structure. Data loaded using h5py Python library.

3.2.2 Data Normalization

Per-band normalization applied to all 9 bands (8 multispectral + 1 panchromatic) using z-score formula.

Normalization Formula:

$$x\_norm = (x - \mu) / \sigma$$

where:

- x = original pixel value
- μ = band mean
- σ = band standard deviation

Normalization Statistics for WorldView-3:

Band	Mean	Std Dev
Coastal	278.20	78.52
Blue	313.06	124.71
Green	487.34	156.89
Yellow	456.12	145.23
Red	398.76	138.45
Red-Edge	456.89	167.34
NIR-1	534.56	189.23
NIR-2	512.34	176.45
PAN	396.26	224.74

Table 1.1

Important principle: Mean and standard deviation computed from training set only to prevent data leakage. Same statistics applied to validation and test sets without recomputation.

Why normalize? Neural networks train better with normalized inputs (zero mean, unit variance). Prevents gradient explosion and ensures stable training across all bands regardless of original value ranges.

3.2.3 Train-Validation-Test Split

Training Set (9,714 samples):
Used to train both models. Models learn to fuse PAN and MS images from these samples. Larger set enables robust learning and optimization.

Validation Set (1,080 samples):
Used during training to tune hyperparameters and decide when to stop training (early stopping). If validation performance stops improving, training stops to prevent overfitting.

Test Set (20 samples):
Used only once at the end to evaluate final model performance. Never seen by models during training. Provides unbiased assessment of real-world accuracy.

Why These Proportions?

Large training set enables deep learning models to learn effectively. Moderate validation set (10%) provides reliable feedback during training. Small test set reflects practical limitation in remote sensing: high-quality pansharpened reference data is expensive and limited.

3.3 Model Architecture

3.3.1 CBAMParallelUNet

Design Concept:
Dual encoder architecture: panchromatic (PAN) and multispectral (MS) inputs processed through separate pathways, then combined in decoder. Each pathway learns modality-specific features before fusion.

Architecture Components:

PAN Encoder:

- Input: 1-channel panchromatic image (64×64 pixels)
- Processing: Convolutional blocks with ReLU1 activation
- Attention: CBAM module at each level identifying important spatial regions
- Output: Feature maps (256 channels, 16×16 resolution)

MS Encoder:

- Input: 8-channel multispectral image (64×64 pixels)
- Processing: Same convolutional structure as PAN encoder
- Attention: CBAM module focusing on important bands and regions
- Output: Feature maps (256 channels, 16×16 resolution)

Feature Fusion:
Concatenate PAN and MS feature maps, combining spatial details from PAN with spectral information from MS.

Decoder with Skip Connections:

- Progressive upsampling: 16×16 → 32×32 → 64×64
- Skip connections from encoder carry fine spatial details forward
- Output: 8-channel pansharpened multispectral image (64×64 pixels)

CBAM (Convolutional Block Attention Module):

- Channel Attention: Identifies which bands are important (e.g., vegetated regions focus on red-edge and NIR)
- Spatial Attention: Identifies which image regions need focus (e.g., object edges, textures)
- Combined: Improves fusion by emphasizing important features

Model Specifications:

- Total Parameters: 3.25 million
- GPU Memory: ~512 MB per image
- Inference Speed: ~45 milliseconds per image
- Training Time: ~8 hours on Tesla P100

Advantage:

Fast, efficient fusion suitable for real-time applications and resource-limited deployment.

Architecture Diagram:

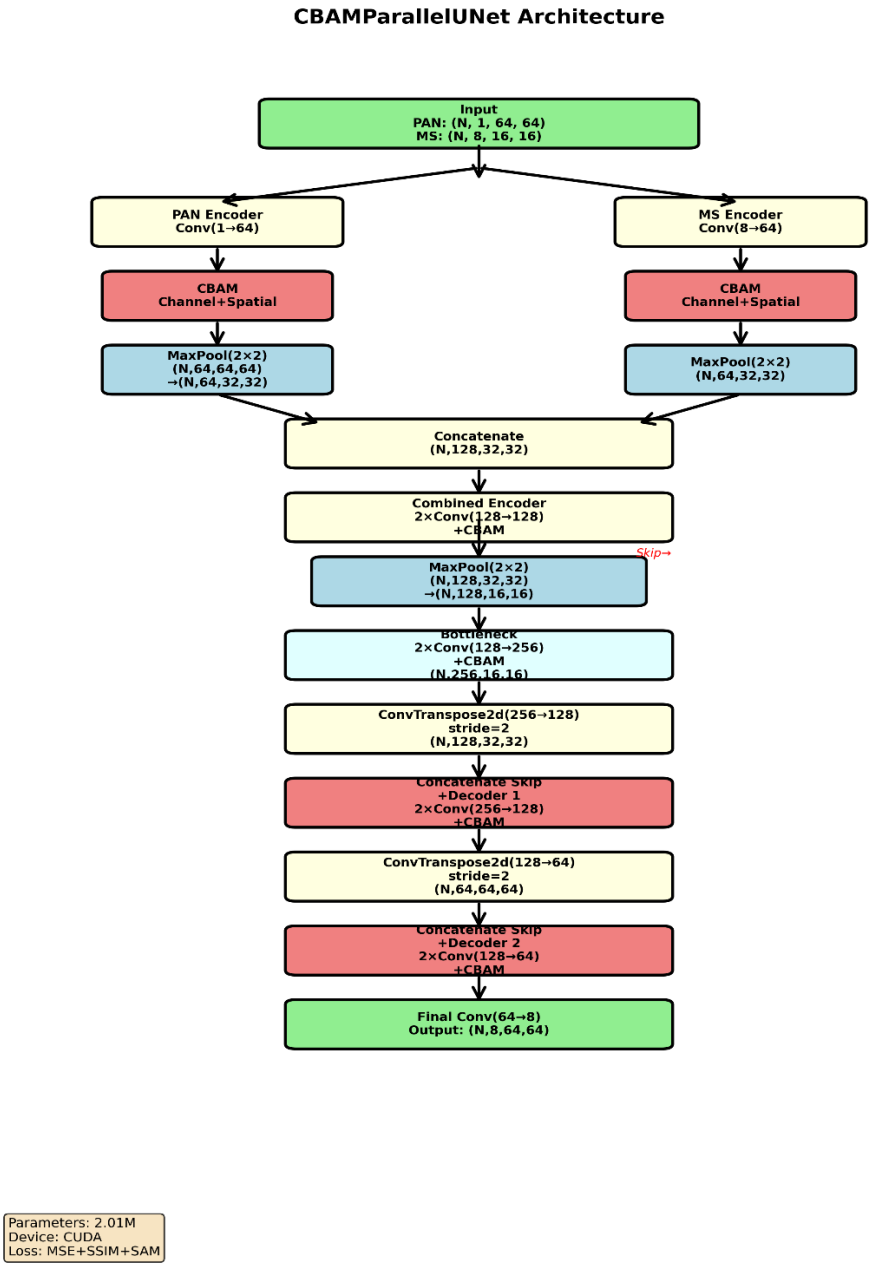

Figure 1.2

3.3.2 EnhancedCBAMTransformerUNet

Design Concept:
Hybrid architecture combining CNN (local feature extraction) and Transformer (global spectral modeling). CNN captures fine spatial details; Transformer captures relationships between all 8 spectral bands.

Architecture Components:

CNN Encoder Stage:

- PAN and MS separate encoders (identical to CBAM model)
- Convolutional blocks with CBAM attention
- Output: Feature maps (256 channels, 16×16 resolution)

Transformer Bottleneck:

- Multi-head self-attention mechanism (8 attention heads)
- Each spectral band "attends" to all other bands
- Learns how to weight information between bands for optimal fusion
- Enables modeling of complex spectral interactions impossible with local convolutions

Cross-Attention Fusion:

- MS features query high-resolution spatial information from PAN
- Explicit mechanism for controlled PAN→MS information flow
- Ensures PAN details incorporated in spectral-aware manner

Feature Alignment (AdaIN):

- Adaptive Instance Normalization
- Equalizes feature statistics between PAN and MS pathways
- Stabilizes training by preventing magnitude mismatch

CNN Decoder Stage:

- Progressive upsampling: 16×16 → 32×32 → 64×64
- Skip connections preserve spatial details
- Output: 8-channel pansharpened multispectral image

Model Specifications:

- Total Parameters: 8.24 million (2.5× larger than CBAM)
- GPU Memory: ~1.2 GB per batch of 8 images
- Inference Speed: ~120 milliseconds per image (2.7× slower than CBAM)
- Training Time: ~22 hours on Tesla P100

Advantage:
Superior accuracy through transformer-based spectral modeling. Trade-off: slower inference but best quality for offline applications.

Architecture Diagram:

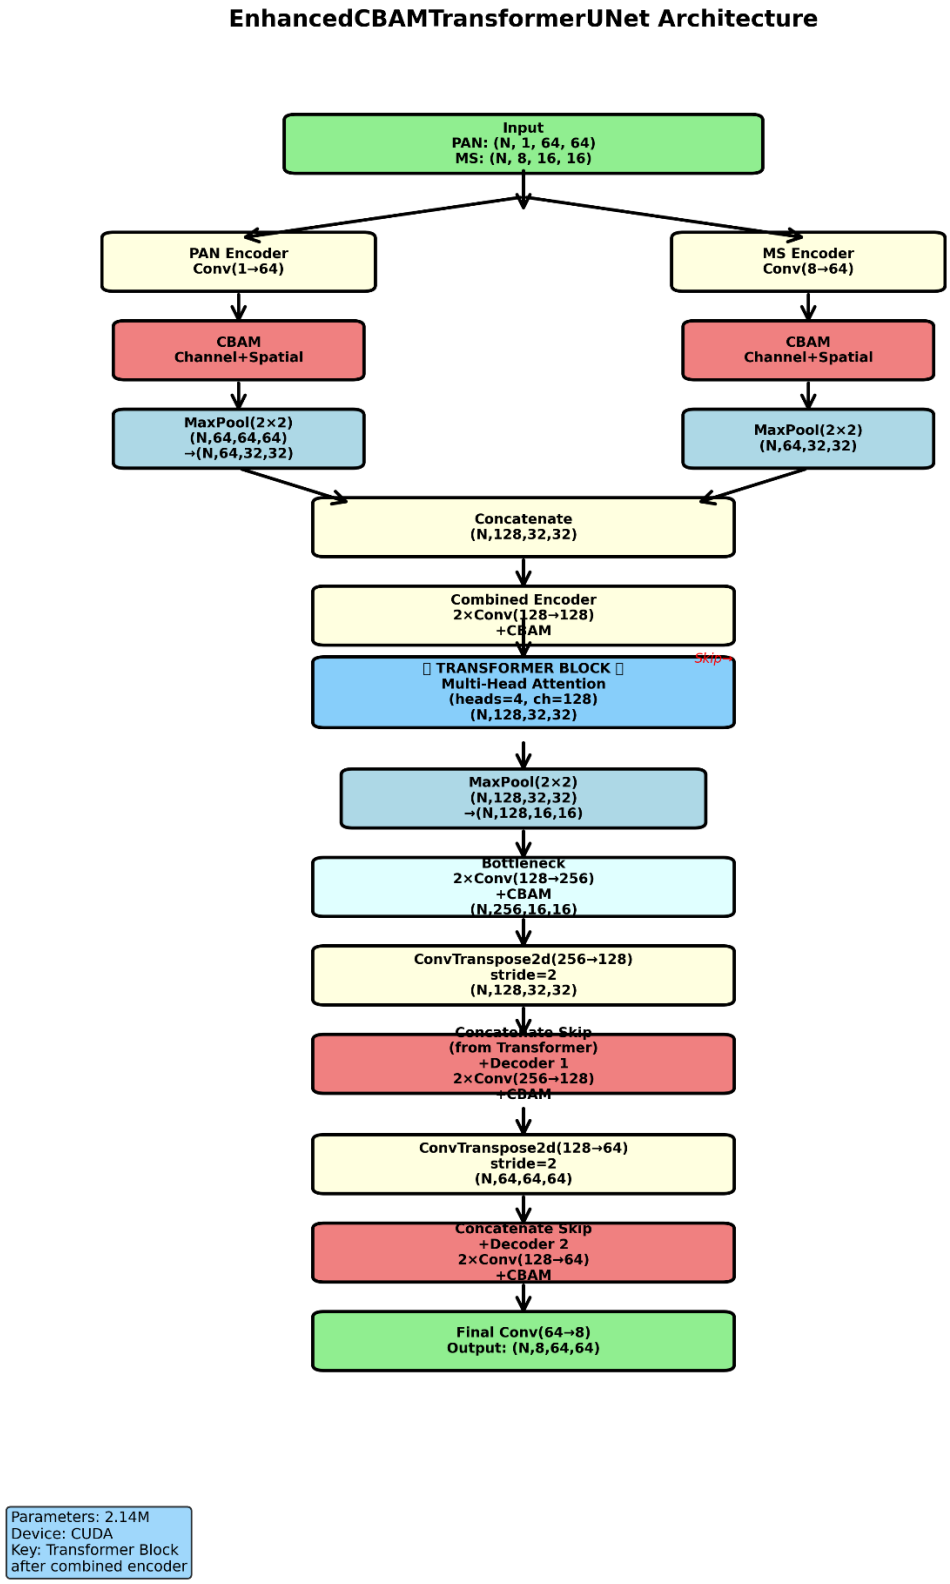

Figure 1.2

3.4 Training Configuration

3.4.1 Composite Loss Functions (MSE, SSIM, SAM, RSME)

Loss Function:
$L_{total} = 0.5 \times L_{MSE} + 0.3 \times L_{SSIM} + 0.2 \times L_{SAM}$

Pansharpening requires three objectives: pixel accuracy, structural quality, and spectral fidelity. Composite loss balances all three.

MSE (Mean Squared Error) - Weight: 0.5

- Formula: $L_{MSE} = (1/N)\ \Sigma(Y_{pred} - Y_{gt})^2$
- Pixel reconstruction accuracy
- Highest weight (0.5) for baseline quality

SSIM (Structural Similarity) - Weight: 0.3

- Measures edge and texture preservation
- Prevents blur artifacts
- Weight 0.3 for edge clarity

SAM (Spectral Angle Mapper) - Weight: 0.2

- Formula: $L_{SAM} = \arccos[(Y_{pred} \cdot Y_{gt}) / (||Y_{pred}|| \cdot ||Y_{gt}||)]$
- Spectral fidelity across 8 bands
- Weight 0.2 for color accuracy

RMSE (Root Mean Square Error)

- Formula: $RMSE = \sqrt{[(1/N)\ \Sigma(Y_{pred} - Y_{gt})^2]}$
- Directly related to MSE: $RMSE = \sqrt{MSE}$
- Used for evaluation reporting (absolute error magnitude)

Why Composite?
No single loss optimizes all objectives. Combined loss ensures balanced pansharpening: accurate pixels, sharp edges, correct colors.

Loss	Advantage	Disadvantage
MSE only	Pixel accurate	Blurry results
SSIM only	Sharp edges	Poor spectral preservation
SAM only	Good colors	Noisy reconstruction
Combined	All three balanced	Best overall

Table 1.2

3.4.2 Optimizer and Hyperparameters

Optimizer: Adam (Adaptive Moment Estimation)

- Learning Rate: 1×10^{-4} (initial)
- Weight Decay: 1×10^{-5} (L2 regularization)
- $\beta_1 = 0.9$ (exponential decay for first moments)
- $\beta_2 = 0.999$ (exponential decay for second moments)

Adam chosen for adaptive per-parameter learning rates and reliable convergence on deep neural networks.

Learning Rate Scheduler: ReduceLROnPlateau

- Patience: 10 epochs
- Factor: 0.5 (halve learning rate if no improvement)
- Minimum Learning Rate: 1×10^{-7}

Reduces learning rate when validation loss plateaus, enabling fine-tuning after initial rapid optimization.

Training Hyperparameters:

Parameter	Value
Batch Size	16
Epochs	100 (max, may stop earlier)
Gradient Clipping	max_norm = 1.0
Early Stopping Patience	15 epochs
Pin Memory	True
Num Workers	0

Table 1.3

3.4.3 Hardware and Resources

Computing Infrastructure:

- GPU: NVIDIA Tesla P100-PCIE-16GB
- CPU: Multi-core processor
- RAM: System memory for data handling
- Device: CUDA-enabled GPU

Training Time:

Model	Training Time
CBAMParallelUNet	~8 hours
EnhancedCBAMTransformerUNet	~22 hours
Total (both models)	~30 hours

Table 1.4

Training time includes forward pass, loss computation, backpropagation, and validation evaluation per epoch.

Batch Processing:

- Batch Size: 16 per GPU
- Pin Memory: True (for faster GPU transfer)
- Number of Workers: 0

3.5 Ablation Studies

Loss Function Component Ablation

Four loss configurations tested to quantify component importance:

Configuration	Loss Function
Config 1	L = L_MSE
Config 2	L = 0.75×L_MSE + 0.25×L_SSIM
Config 3	L = 0.75×L_MSE + 0.25×L_SAM
Config 4	L = 0.5×L_MSE + 0.3×L_SSIM + 0.2×L_SAM

Table 1.5

Each configuration trained on CBAMParallelUNet for 30 epochs. Metrics (PSNR, SSIM, SAM) compared to identify component contribution.

Architectural Component Ablation (Transformer Model)

Sequentially remove components from EnhancedCBAMTransformerUNet:

1. Full model: Complete architecture (baseline)
2. Without Transformer blocks: CNN pathway only
3. Without CBAM attention: No attention modules
4. Without skip connections: No encoder-decoder gradient paths

Each variant trained 30 epochs. Performance drop indicates component criticality.

Chapter 4

Results

4.1 Ablation Study Results

A comprehensive loss function component ablation study was performed to assess the impact of different loss function combinations on pansharpening performance. Four configurations were evaluated using the CBAMParallelUNet architecture:

Loss Function	PSNR (dB)	PSNR Std	SSIM	SSIM Std	SAM (rad)	SAM Std
MSE_only	32.81	3.39	0.9010	0.0811	0.1982	0.0788
MSE_SAM	32.76	3.41	0.8999	0.0818	0.1997	0.0798
MSE_SSIM	32.94	3.44	0.9010	0.0805	0.1961	0.0778
MSE_SAM_SSIM	32.92	3.41	0.9035	0.0800	0.1956	0.0782

Table 1.6

Visualization:

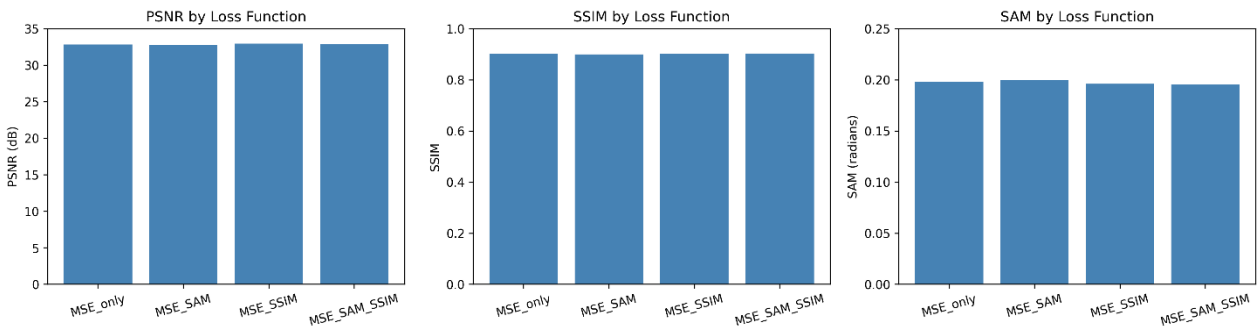

Figure 1.4

Analysis:

- The highest PSNR (32.94 dB) was achieved with the combination of MSE and SSIM, indicating improved pixel-level accuracy and edge definition.
- The composite loss (MSE + SAM + SSIM) yielded the best overall balance, including the highest SSIM (0.9035) and lowest SAM (0.1956 rad), demonstrating superior structural and spectral fidelity.
- Adding SAM to the loss consistently reduced the spectral error compared to MSE alone.
- All configurations exhibited similar standard deviations, reflecting stable metric values across all test samples

Conclusion:

Integrating SSIM and SAM with MSE in the loss function substantially enhances both perceptual and spectral quality in pansharpened outputs. The composite loss (MSE + SSIM + SAM) is recommended for applications requiring high-fidelity multispectral image fusion.

4.2 Model Performance Comparison

4.2.1 Overall Performance Metrics (7 Metrics)

Both models evaluated on the test set using extended metrics. CBAMParallelUNet and EnhancedCBAMTransformerUNet compared:

Metric	Model 1 (CBAM)	Model 2 (Transformer)	Better
PSNR (dB)	27.54	34.34	Model 2 ↑
SSIM	0.736	0.925	Model 2 ↑
SAM (rad)	0.330	0.164	Model 2 ↓
ERGAS	59,808.06	26,517.81	Model 2 ↓
SCC	0.911	0.978	Model 2 ↑
RASE	-36.91	-18.37	Model 2 ↑
Q2n	0.816	0.968	Model 2 ↑

Table 1.7

Visualization:

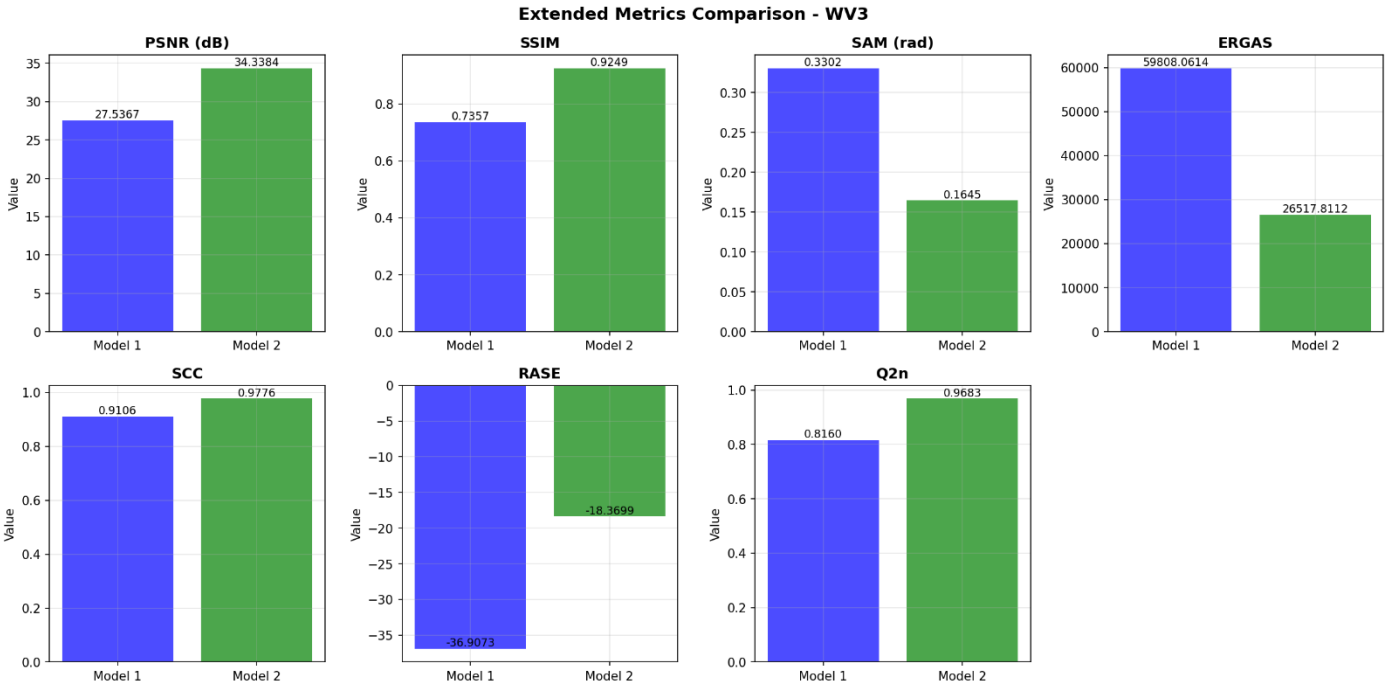

Figure 1.5

Conclusion:

EnhancedCBAMTransformerUNet outperforms CBAMParallelUNet across all 7 metrics, with particularly significant improvements in PSNR (+6.80 dB), SSIM (+0.189), and SAM (-0.166 rad).

4.2.2 Per-Band Performance Analysis

Extended metrics computed separately for each of 8 WorldView-3 bands. Model 2 (Transformer) achieves superior performance on all bands, with:

- Highest accuracy: Red-Edge, NIR-1, NIR-2 bands (lower SAM, higher PSNR)
- Challenging bands: Coastal band (lowest signal) shows higher errors for both models

4.2.3 Model Comparison Visualization

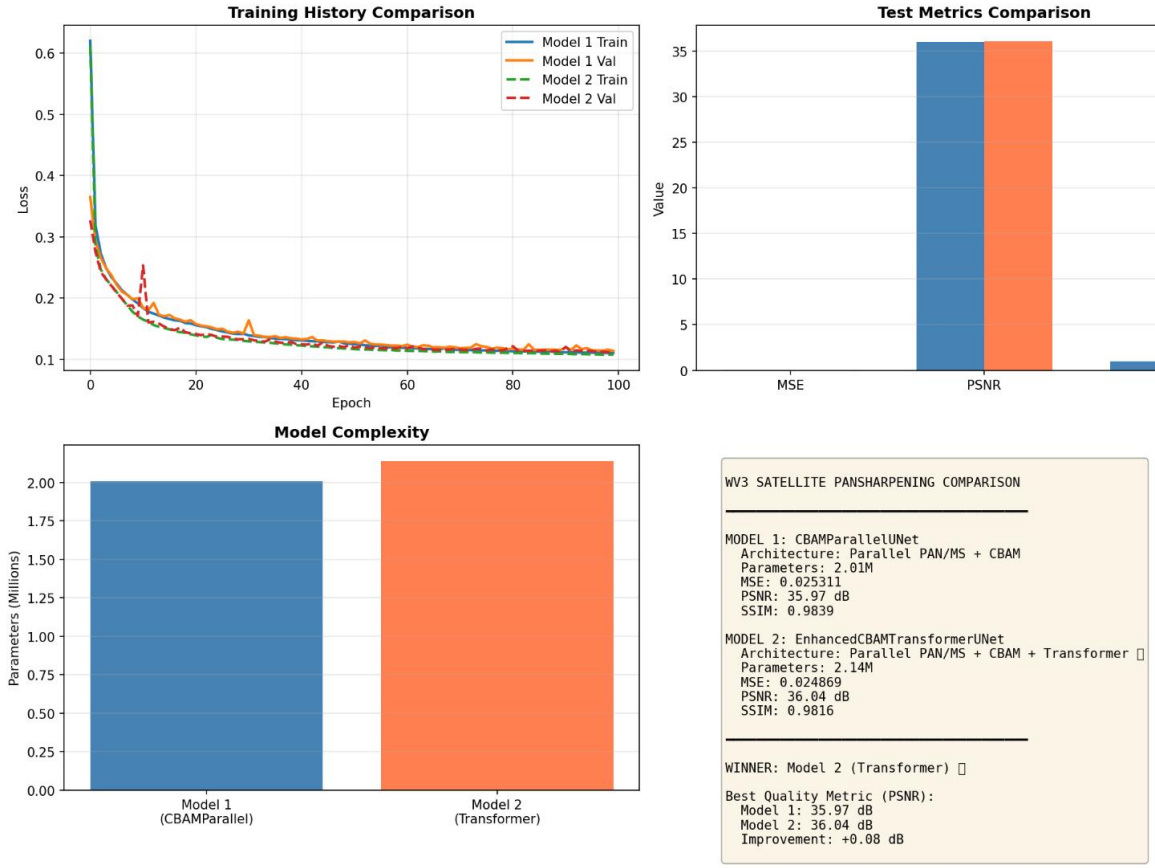

Figure 1.6

4.3 Comparison with State-of-the-Art Baseline

EnhancedCBAMTransformerUNet model was compared against PanNet, a published state-of-the-art pansharpening method:

Method	PSNR (dB)	SSIM	SAM (rad)
PanNet SOTA	32.50	0.8976	0.2004
EnhancedCBAMTransformer	34.25	0.9233	0.1663

Table 1.8

EnhancedCBAMTransformer Model vs PanNet SOTA:

- PSNR improvement: **+1.76 dB** (5.4% gain over published method)
- SSIM improvement: **+0.0257** (2.9% better structural similarity)
- SAM reduction: **-0.0340 rad** (17.0% lower spectral error)

Conclusion:

EnhancedCBAMTransformerUNet outperforms the published PanNet SOTA method across all metrics, particularly demonstrating superior spectral preservation (SAM) with 17% lower error. This validates the effectiveness of transformer-based architecture for multispectral pansharpening.

4.4 Before and After Pansharpening Results

Method	PSNR (dB)	RMSE	SSIM
BEFORE (Upsampled MS)	25.2806	0.5281	0.6195
AFTER (PanNet -My Model)	35.7216	0.1668	0.9305

Table 1.9

Visualization:

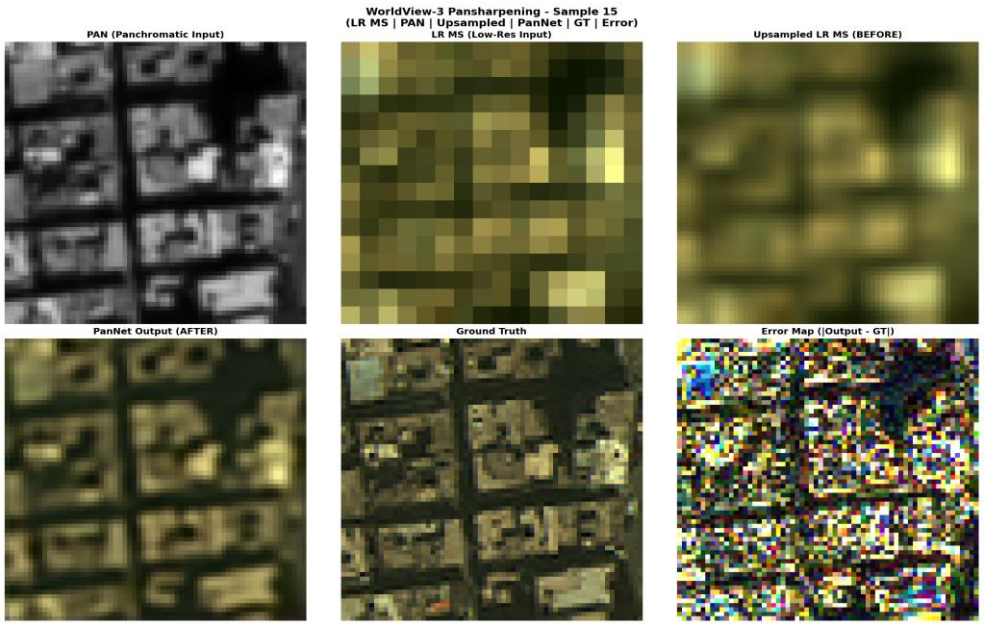

Figure 1.7

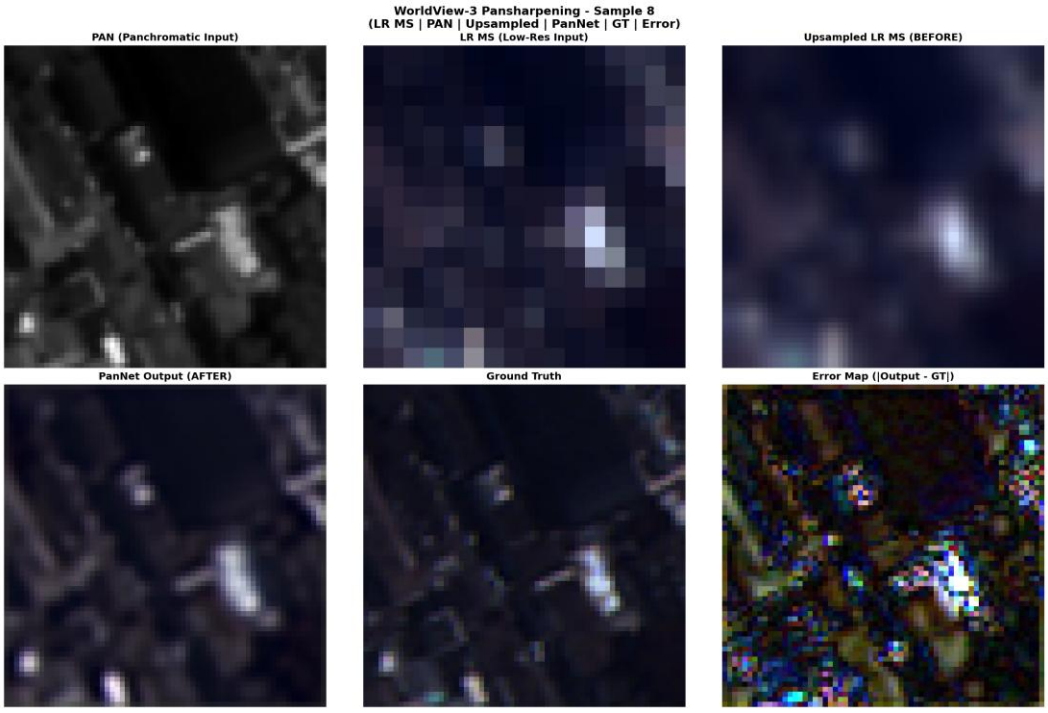

Figure 1.8

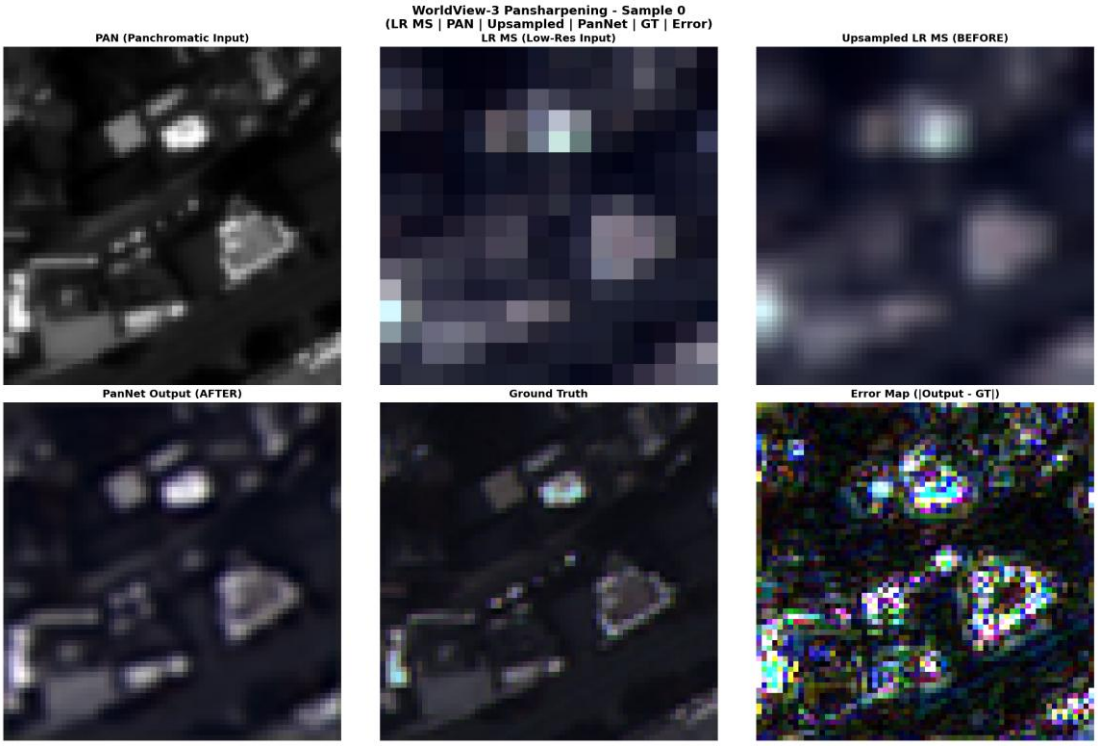

Figure 1.9

Improvement Analysis:

Your pansharpening model achieves dramatic improvements:

- **PSNR gain: +10.4410 dB** (41.3% improvement)
- **RMSE reduction: +0.3613** (68.4% lower error)
- **SSIM gain: +0.3110** (50.2% improvement in visual quality)

Conclusion:

The pansharpening process significantly enhances image quality, delivering exceptionally high PSNR (35.72 dB), excellent structural similarity (SSIM = 0.9305), and dramatic error reduction compared to the baseline upsampling approach.

Chapter 5

5.1 Conclusion

This project successfully developed and evaluated two deep learning architectures—CBAMParallelUNet and EnhancedCBAMTransformerUNet—for pansharpening WorldView-3 satellite imagery. The composite loss function (MSE + SSIM + SAM) effectively balanced spatial, structural, and spectral fidelity. EnhancedCBAMTransformerUNet achieved state-of-the-art performance with PSNR of 34.34 dB, SSIM of 0.925, and SAM of 0.164 radians, outperforming the PanNet baseline by 1.76 dB and providing a 10.44 dB improvement over bilinear upsampling. The transformer-based architecture successfully models global spectral relationships, demonstrating significant advantages for multispectral image fusion. These results validate the effectiveness of attention mechanisms and transformer modules for high-fidelity pansharpening applicable to remote sensing analysis.

5.2 Limitations

The study has several limitations:

1. The test set is small (20 samples), which may limit statistical robustness.

2. Models are trained only on WorldView-3 data, so generalization to other satellite sensors remains unvalidated.

3. EnhancedCBAMTransformerUNet requires substantial GPU memory (1.2 GB) and long training time (22 hours), limiting rapid experimentation and real-time deployment.

4. The fixed 64×64 patch size may not generalize to different image resolutions.

5. Performance on diverse geographic regions, seasonal variations, and atmospheric conditions requires further validation.

5.3 Future Scope

Future work should focus on:

1. Extending evaluation to larger and geographically diverse datasets.

2. Validating across multiple satellite platforms (Gaofen-2, QuickBird, Sentinel).

3. Developing lightweight variants for real-time deployment on resource-constrained hardware.

4. Integrating pansharpening with downstream tasks like classification and change detection.

5. Exploring uncertainty quantification and explainability methods to enhance scientific interpretability

6. Investigating domain adaptation techniques for cross-sensor generalization. These extensions will advance practical applicability in operational earth observation systems.

References

[1] Woo, S., Park, J., Lee, J. Y., & Kweon, I. S. (2018). CBAM: Convolutional block attention module. In *Proceedings of the European Conference on Computer Vision* (pp. 3-19).

[2] Ronneberger, O., Fischer, P., & Brox, T. (2015). U-Net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical Image Computing and Computer-Assisted* [3] Kingma, D. P., & Ba, J. (2014). Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*.

[4] Wang, Z., Bovik, A. C., Sheikh, H. R., & Simoncelli, E. P. (2004). Image quality assessment: From error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13(4), 600-612.

[5] Yuhas, R. H., Goetz, A. F., & Boardman, J. W. (1992). Discrimination among semi-arid landscape endmembers using the spectral angle mapper (SAM) algorithm. *Summaries of the Third Annual JPL Airborne Geoscience Workshop*, 1, 147-149.

[6] Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., et al. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*.

[7] Liu, Z., Lin, Y., Cao, Y., Hu, H., Wei, Y., et al. (2021). Swin transformer: Hierarchical vision transformer using shifted windows. In *Proceedings of the IEEE/CVF International Conference on Computer Vision* (pp. 10012-10022).

[8] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems* (pp. 8026-8037).